

# Regression Discontinuity

Alex Alekseev

```
library(tidyverse)
library(rdrobust)
library(rddensity)
```

## Intro

In this session, we will practice using the regression discontinuity design (RDD) using a simulated dataset. There are several packages in R that implement RDD, we will be using `rdrobust` and `rddensity`.

## Data

```
nobs <- 1000
thresh <- 0.5
jump <- 0.25
```


We will generate a dataset with 1000 observations of the running variable ( $x$ ), potential outcomes ( $y_0, y_1$ ), observed outcomes ( $y$ ), and the treatment status ( $d$ ). We will also create a centered running variable ( $x\_cent$ ) for which the threshold is 0.

```
set.seed(123)

df <- tibble(
  x = rnorm(n = nobs)
  , y0 = x^2/100
  , y1 = y0 + jump
  , d = 1*(x >= thresh)
  , y = (1 - d)*y0 + d*y1 + rnorm(n = nobs, sd = 0.3)
  , x_cent = x - thresh
)
```

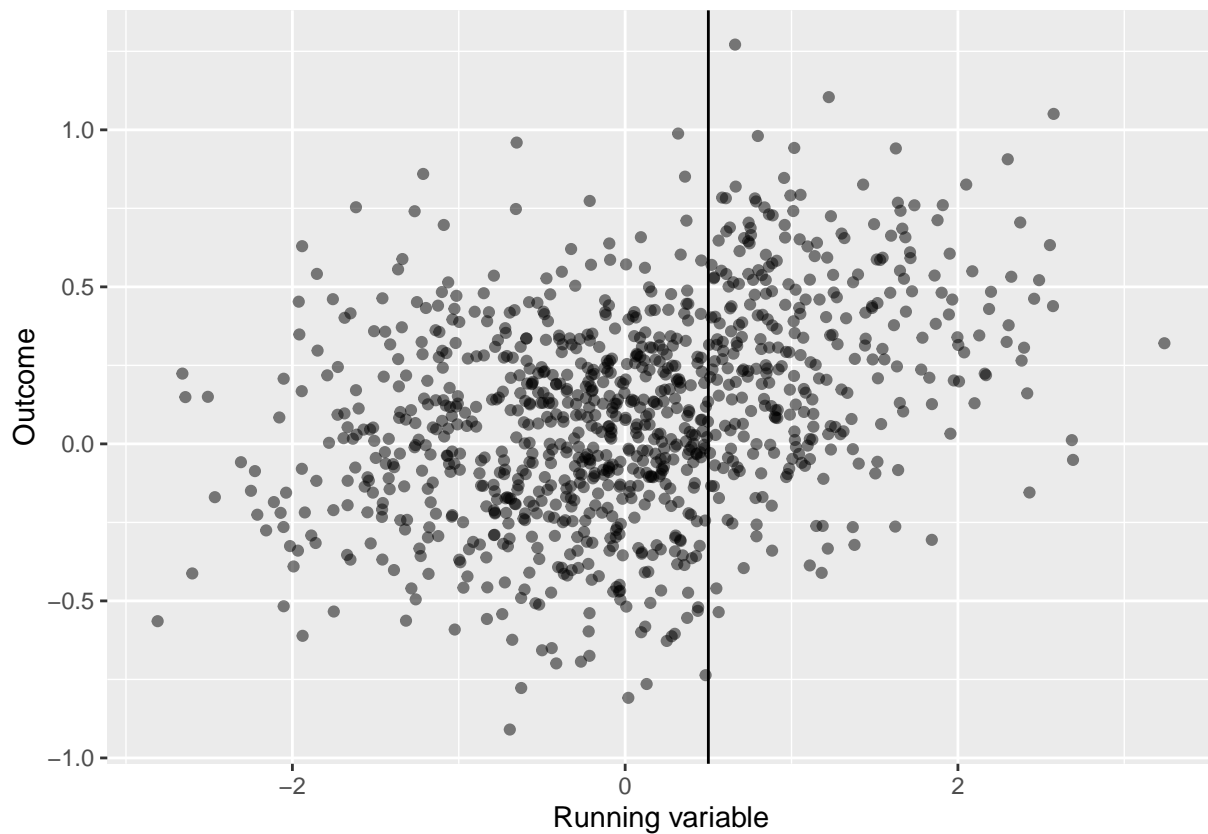
The potential outcome  $Y^0$  is a non-linear (quadratic) function of the running variable. The potential outcome  $Y^1$  is simply shifted by fixed amount (0.25) relative to  $Y^0$ . This set-up satisfies the main identifying assumption of RDD: *continuity*. The only reason why you would observe a jump in the observed outcomes is due to treatment. In the absence of treatment, the potential outcomes are continuous at the threshold (0.5).

```
df %>%
  pivot_longer(
    cols = c(y0, y1)
    , names_to = "Treatment_state"
    , values_to = "Potential_outcomes"
  ) %>%
  ggplot(aes(x, Potential_outcomes, color = Treatment_state)) +
  geom_point(alpha = 0.3) +
  geom_vline(xintercept = 0.5) +
  xlab("Running variable")
```




Let's switch from the potential outcomes to the observed outcomes.

```
ggplot(df, aes(x, y)) +
  geom_point(alpha = 0.5) +
  geom_vline(xintercept = 0.5) +
  labs(x = "Running variable", y = "Outcome")
```



There is a lot of noise in the data, which makes it hard to see a jump at the threshold. We can add trend lines on both sides of the cutoff to aid the analysis.

```
ggplot(df, aes(x, y, color = (x >= thresh))) +  
  geom_point(alpha = 0.5) +  
  geom_smooth(method = "loess", se = F) +  
  geom_vline(xintercept = 0.5) +  
  labs(x = "Running variable", y = "Outcome")
```



Now the jump at the threshold becomes more pronounced. We can visually confirm that this jump is in fact around the true value of 0.5.

## Estimation

To estimate the treatment effect formally, we will use the `rdrobust` function.

```
m <- rdrobust(
  df$y, df$x_cent, c = 0
  , all = T
)
summary(m)
```


```
## Sharp RD estimates using local polynomial regression.
##
## Number of Obs.          1000
## BW type                mserd
## Kernel                  Triangular
## VCE method              NN
##
## Number of Obs.          697      303
## Eff. Number of Obs.     185      133
## Order est. (p)          1         1
## Order bias (q)          2         2
## BW est. (h)             0.448     0.448
## BW bias (b)             0.707     0.707
```

```
## rho (h/b)                0.633        0.633
## Unique Obs.              697          303
##
## =====
##          Method      Coef. Std. Err.      z    P>|z|      [ 95% C.I. ]
## =====
## Conventional      0.204      0.075      2.698    0.007    [0.056 , 0.351]
## Bias-Corrected    0.185      0.075      2.453    0.014    [0.037 , 0.333]
## Robust            0.185      0.090      2.064    0.039    [0.009 , 0.361]
## =====
```

The option `all = T` allows us to see different estimates (conventional or bias-corrected) and standard errors (conventional or robust). The estimated effect is between 0.204 and 0.185, depending on the estimate, which is close to the true value of 0.25 albeit lower. The confidence intervals do include the true value.

We can use the `rdplot` function to visualize the results.

```
rdplot(df$y, df$x_cent)
```



Notice that the function made some choices for us. For example, it picked the number of polynomials  $p$  ( $= 1$ ) in the local regression and the bandwidth  $h$  ( $= 0.448$ ). Let's change these numbers manually and see if it changes the results.

First, let's increase the degree of polynomials by changing the value of  $p$ .

```
m1 <- rdrobust(
  df$y, df$x_cent, c = 0
  , all = T
```

```
, p = 2
)
summary(m1)
```

```
## Sharp RD estimates using local polynomial regression.
```

```
##
```

```
## Number of Obs.          1000
```

```
## BW type                  mserd
```

```
## Kernel                   Triangular
```

```
## VCE method               NN
```

```
##
```

```
## Number of Obs.          697          303
```

```
## Eff. Number of Obs.     255          172
```

```
## Order est. (p)          2            2
```

```
## Order bias (q)          3            3
```

```
## BW est. (h)             0.609        0.609
```

```
## BW bias (b)             0.862        0.862
```

```
## rho (h/b)              0.707        0.707
```

```
## Unique Obs.            697          303
```

```
##
```

```
## =====
```

```
##           Method      Coef. Std. Err.      z    P>|z|      [ 95% C.I. ]
```

```
## =====
```

```
##   Conventional      0.200    0.093    2.140    0.032    [0.017 , 0.383]
```

```
## Bias-Corrected      0.192    0.093    2.057    0.040    [0.009 , 0.375]
```

```
##           Robust      0.192    0.106    1.817    0.069    [-0.015 , 0.399]
```

```
## =====
```

Increasing p widened the confidence intervals.

This pattern continues if we further increase p.

```
m2 <- rdrobust(
  df$y, df$x_cent, c = 0
  , all = T
  , p = 4
)
summary(m2)
```

```
## Sharp RD estimates using local polynomial regression.
```

```
##
```

```
## Number of Obs.          1000
```

```
## BW type                  mserd
```

```
## Kernel                   Triangular
```

```
## VCE method               NN
```

```
##
```

```
## Number of Obs.          697          303
```

```
## Eff. Number of Obs.     360          217
```

```
## Order est. (p)          4            4
```

```
## Order bias (q)          5            5
```

```
## BW est. (h)          0.878          0.878
## BW bias (b)          1.128          1.128
## rho (h/b)            0.778          0.778
## Unique Obs.          697           303
##
## =====
##      Method      Coef. Std. Err.      z    P>|z|      [ 95% C.I. ]
## =====
##   Conventional    0.213    0.127    1.682    0.093    [-0.035 , 0.461]
## Bias-Corrected    0.220    0.127    1.739    0.082    [-0.028 , 0.468]
##      Robust       0.220    0.136    1.623    0.105    [-0.046 , 0.486]
## =====
```

Increasing  $p$  beyond the default value seems to produce worse results.

Let's set the bandwidth  $h$  manually. We will start with a value smaller than the computed default value.

```
m3 <- rdrobust(
  df$y, df$x_cent, c = 0
  , all = T
  , h = 0.5
)
summary(m3)
```

```
## Sharp RD estimates using local polynomial regression.
##
## Number of Obs.          1000
## BW type                 Manual
## Kernel                  Triangular
## VCE method              NN
##
## Number of Obs.          697           303
## Eff. Number of Obs.     202           146
## Order est. (p)           1             1
## Order bias (q)           2             2
## BW est. (h)              0.500         0.500
## BW bias (b)              0.500         0.500
## rho (h/b)                1.000         1.000
## Unique Obs.              697           303
##
## =====
##      Method      Coef. Std. Err.      z    P>|z|      [ 95% C.I. ]
## =====
##   Conventional    0.209    0.073    2.879    0.004    [0.067 , 0.351]
## Bias-Corrected    0.182    0.073    2.507    0.012    [0.040 , 0.324]
##      Robust       0.182    0.103    1.767    0.077    [-0.020 , 0.384]
## =====
```

Notice two things. First, there is a drop in the effective number of observations (Eff.

Number of Obs.) relative to the very first specification. Making the bandwidth smaller means we are using fewer observations that are closer to the threshold. Second, this tends to make the estimate less precise (the standard errors go up) but also, interestingly, pushes the estimated effect up.

Let's try to pick a larger bandwidth.

```
m4 <- rdrobust(
  df$y, df$x_cent, c = 0
  , all = T
  , h = 1
)
summary(m4)
```

```
## Sharp RD estimates using local polynomial regression.
```

```
##
## Number of Obs.          1000
## BW type                 Manual
## Kernel                  Triangular
## VCE method              NN
##
## Number of Obs.          697      303
## Eff. Number of Obs.     402      229
## Order est. (p)           1         1
## Order bias (q)           2         2
## BW est. (h)              1.000     1.000
## BW bias (b)              1.000     1.000
## rho (h/b)                1.000     1.000
## Unique Obs.              697      303
##
## =====
##           Method      Coef. Std. Err.      z    P>|z|      [ 95% C.I. ]
## =====
##   Conventional      0.235    0.055    4.291    0.000    [0.127 , 0.342]
## Bias-Corrected      0.200    0.055    3.664    0.000    [0.093 , 0.308]
##           Robust      0.200    0.077    2.594    0.009    [0.049 , 0.352]
## =====
```

Larger bandwidth increases the effective number of observations and makes the estimates more precise.

We can conduct a sensitivity analysis by showing how our estimates would change with the bandwidth.

```
est_fun <- function(x) {
  m <- rdrobust(
    df$y, df$x_cent, c = 0
    , all = T
    , h = x
  )
}
```




```

est <- m$Estimate[1, "tau.bc"]
ci_low <- m$ci[3, 1]
ci_high <- m$ci[3, 2]
res <- c("est" = est, "ci_low" = ci_low, "ci_high" = ci_high)
return(res)
}

df_est <- tibble(
  bandwidth = seq(0.2, 2, 0.2)
  , estimate = map_df(bandwidth, est_fun)
) %>%
  unnest(cols = estimate)

ggplot(df_est, aes(x = bandwidth, y = est.tau.bc)) +
  geom_hline(yintercept = jump, linetype = "dashed") +
  geom_errorbar(aes(ymin = ci_low, ymax = ci_high), color = "gray50", width = 0.05) +
  geom_point() +
  geom_line() +
  labs(y = "Effect")

```



We see that the estimates are pretty stable and the confidence intervals include the true value.

## Robustness checks

We can conduct some placebo tests by picking the “wrong” thresholds. As the name suggests, we should not find any effects there.

We find a zero effect at -0.5 (left side of the threshold).

```
m5 <- rdrobust(
  df$y, df$x, c = -0.5
  , all = T
)
summary(m5)
```

```
## Sharp RD estimates using local polynomial regression.
```

```
##
```

```
## Number of Obs.          1000
```

```
## BW type                mserd
```

```
## Kernel                  Triangular
```

```
## VCE method              NN
```

```
##
```

```
## Number of Obs.          295          705
```

```
## Eff. Number of Obs.     136          209
```

```
## Order est. (p)          1            1
```

```
## Order bias (q)          2            2
```

```
## BW est. (h)             0.522        0.522
```

```
## BW bias (b)             0.818        0.818
```

```
## rho (h/b)              0.638        0.638
```

```
## Unique Obs.            295          705
```

```
##
```

```
## =====
```

```
##           Method      Coef. Std. Err.      z    P>|z|      [ 95% C.I. ]
```

```
## =====
```

```
##   Conventional      0.044    0.066    0.667    0.505    [-0.085 , 0.173]
```

```
## Bias-Corrected      0.044    0.066    0.673    0.501    [-0.085 , 0.174]
```

```
##           Robust      0.044    0.078    0.569    0.569    [-0.109 , 0.197]
```

```
## =====
```

We find a zero effect at 1 (right side of the threshold).

```
m6 <- rdrobust(
  df$y, df$x, c = 1
  , all = T
)
summary(m6)
```

```
## Sharp RD estimates using local polynomial regression.
```

```
##
```

```
## Number of Obs.          1000
```

```
## BW type                mserd
```

```
## Kernel                  Triangular
```

```
## VCE method              NN
```

```
##
## Number of Obs.                843        157
## Eff. Number of Obs.          94         67
## Order est. (p)                1          1
## Order bias (q)               2          2
## BW est. (h)                  0.338      0.338
## BW bias (b)                  0.644      0.644
## rho (h/b)                    0.525      0.525
## Unique Obs.                  843        157
##
## =====
##      Method      Coef. Std. Err.      z    P>|z|      [ 95% C.I. ]
## =====
##      Conventional    -0.009    0.109    -0.081    0.935    [-0.223 , 0.205]
##      Bias-Corrected    0.033    0.109    0.306    0.760    [-0.181 , 0.247]
##      Robust          0.033    0.124    0.270    0.787    [-0.209 , 0.276]
## =====
```

We can also conduct the density test to check for manipulation around the threshold.

```
rdd_dens <- rddensity(df$x_cent)
summary(rdd_dens)
```

```
##
## Manipulation testing using local polynomial density estimation.
##
## Number of obs =          1000
## Model =              unrestricted
## Kernel =             triangular
## BW method =          estimated
## VCE method =         jackknife
##
## c = 0                Left of c      Right of c
## Number of obs        697           303
## Eff. Number of obs    271           203
## Order est. (p)        2            2
## Order bias (q)        3            3
## BW est. (h)           0.664         0.758
##
## Method                T              P > |T|
## Robust                -0.2746        0.7836
##
##
## P-values of binomial tests (H0: p=0.5).
##
## Window Length / 2      <c      >=c      P>|T|
## 0.062                  20       20       1.0000
## 0.125                  46       38       0.4452
## 0.187                  69       59       0.4264
```

## 0.249	93	76	0.2183
## 0.311	122	93	0.0559
## 0.374	147	115	0.0553
## 0.436	180	132	0.0077
## 0.498	201	146	0.0037
## 0.560	224	163	0.0022
## 0.623	260	177	0.0001

The  $p$ -value is 0.784, which suggests that we cannot reject the null of no discontinuities in the density of the running variable. We conclude that there is no manipulation around the threshold. We know from our data simulation that this is indeed the case.

We can also present the results visually.

```
rdplotdensity(rdd_dens, df$x_cent)$Estplot
```

